

Libraries and Configurations

```
import numpy as np
import pandas as pd
import tensorflow as tf
import tensorflow.keras.layers as L
import tensorflow_addons as tfa
import glob, random, os, warnings
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns

print('TensorFlow Version ' + tf.__version__)

import warnings
warnings.filterwarnings("ignore")

TensorFlow Version 2.6.4

image_size = 224
batch_size = 16
n_classes = 3
EPOCHS = 30

train_path = '/kaggle/input/pavement-crack-datasetinfrared-only/train'

classes = {1 : "High Crack",
           2 : "Low Crack",
           3 : "Medium Crack",
           4 : "No Crack"}
```

Data Augmentations

```
def data_augment(image):
    p_spatial = tf.random.uniform([], 0, 1.0, dtype = tf.float32)
    p_rotate = tf.random.uniform([], 0, 1.0, dtype = tf.float32)
    p_pixel_1 = tf.random.uniform([], 0, 1.0, dtype = tf.float32)
    p_pixel_2 = tf.random.uniform([], 0, 1.0, dtype = tf.float32)
    p_pixel_3 = tf.random.uniform([], 0, 1.0, dtype = tf.float32)

    # Flips
    image = tf.image.random_flip_left_right(image)
    image = tf.image.random_flip_up_down(image)

    if p_spatial > .75:
        image = tf.image.transpose(image)
```

```

# Rotates
if p_rotate > .75:
    image = tf.image.rot90(image, k = 3) # rotate 270°
elif p_rotate > .5:
    image = tf.image.rot90(image, k = 2) # rotate 180°
elif p_rotate > .25:
    image = tf.image.rot90(image, k = 1) # rotate 90°

# Pixel-level transforms
# if p_pixel_1 >= .4:
#     image = tf.image.random_saturation(image, lower = .7, upper = 1.3)
# if p_pixel_2 >= .4:
#     image = tf.image.random_contrast(image, lower = .8, upper = 1.2)
# if p_pixel_3 >= .4:
#     image = tf.image.random_brightness(image, max_delta = .1)

return image

```

Data Generator

```

datagen = tf.keras.preprocessing.image.ImageDataGenerator(rescale = 1./255,

samplewise_center = True,

samplewise_std_normalization = True,

preprocessing_function = data_augment)

# set as training data

train_datagen = tf.keras.preprocessing.image.ImageDataGenerator(
    rescale=1.0 / 255,
    rotation_range=30,
    width_shift_range=0.3,
    height_shift_range=0.3,
    shear_range=0.2,
    zoom_range=0.3,
    horizontal_flip=True,
    vertical_flip=True,
)

test_datagen =
tf.keras.preprocessing.image.ImageDataGenerator(rescale=1./255)

```

```

train_gen = datagen.flow_from_directory(
    train_path,
    target_size=(224, 224),
    batch_size = batch_size,
    seed = 1,
    color_mode = 'rgb',
    shuffle = True,
    class_mode='categorical',
)

# same directory as training data

valid_gen = test_datagen.flow_from_directory(
    "/kaggle/input/pavement-crack-dataset/infrared-only/test",
    target_size=(224, 224),
    batch_size = batch_size,
    seed = 1,
    color_mode = 'rgb',
    shuffle = False,
    class_mode='categorical'
)

```

Found 1855 images belonging to 4 classes.
Found 461 images belonging to 4 classes.

Sample Image Visualization

```

warnings.filterwarnings("ignore")


images = [train_gen[0][0][i] for i in range(16)]
fig, axes = plt.subplots(3, 5, figsize = (10, 10))

axes = axes.flatten()

for img, ax in zip(images, axes):
    ax.imshow(img.reshape(image_size, image_size, 3))
    ax.axis('off')

plt.tight_layout()
plt.show()

```



Building the Model

```
!pip install vit_keras
Collecting vit_keras
  Downloading vit_keras-0.1.2-py3-none-any.whl (24 kB)
Collecting validators
  Downloading validators-0.20.0.tar.gz (30 kB)
  Preparing metadata (setup.py) ... ent already satisfied: scipy in
/opt/conda/lib/python3.7/site-packages (from vit_keras) (1.7.3)
Requirement already satisfied: numpy<1.23.0,>=1.16.5 in
/opt/conda/lib/python3.7/site-packages (from scipy->vit_keras)
(1.21.6)
Requirement already satisfied: decorator>=3.4.0 in
/opt/conda/lib/python3.7/site-packages (from validators->vit_keras)
(5.1.1)
Building wheels for collected packages: validators
  Building wheel for validators (setup.py) ... e=validators-0.20.0-
```

```
py3-none-any.whl size=19582
sha256=6cff2e8546df2ee65d3ed6aaa9d25eba9b1f516c5704bcd6f12109e21935e71
1
  Stored in directory:
/root/.cache/pip/wheels/5f/55/ab/36a76989f7f88d9ca7b1f68da6d94252bb6a8
d6ad4f18e04e9
Successfully built validators
Installing collected packages: validators, vit_keras
Successfully installed validators-0.20.0 vit_keras-0.1.2
WARNING: Running pip as the 'root' user can result in broken
permissions and conflicting behaviour with the system package manager.
It is recommended to use a virtual environment instead:
https://pip.pypa.io/warnings/venv
```

1 - ViT B16 Model

```
from vit_keras import vit

vit_model = vit.vit_b16(
    image_size = image_size,
    activation = 'softmax',
    pretrained = True,
    include_top = False,
    pretrained_top = False,
    classes = 4)

Downloading data from
https://github.com/faustomoraes/vit-keras/releases/download/dl/ViT-
B_16_imagenet21k+imagenet2012.npz
347504640/347502902 [=====] - 3s 0us/step
347512832/347502902 [=====] - 3s 0us/step

class Patches(L.Layer):
    def __init__(self, patch_size):
        super(Patches, self).__init__()
        self.patch_size = patch_size

    def call(self, images):
        batch_size = tf.shape(images)[0]
        patches = tf.image.extract_patches(
            images = images,
            sizes = [1, self.patch_size, self.patch_size, 1],
            strides = [1, self.patch_size, self.patch_size, 1],
            rates = [1, 1, 1, 1],
            padding = 'VALID',
        )
        patch_dims = patches.shape[-1]
```

```
patches = tf.reshape(patches, [batch_size, -1, patch_dims])
return patches
```

Visualizing Attention Maps of Sample Test Image

```
plt.figure(figsize=(4, 4))
batch_size = 16
patch_size = 7 # Size of the patches to be extract from the input images
num_patches = (image_size // patch_size) ** 2

x = train_gen.next()
image = x[0][0]

plt.imshow(image.astype('uint8'))
plt.axis('off')



resized_image = tf.image.resize(
    tf.convert_to_tensor([image]), size = (image_size, image_size)
)

patches = Patches(patch_size)(resized_image)
print(f'Image size: {image_size} X {image_size}')
print(f'Patch size: {patch_size} X {patch_size}')
print(f'Patches per image: {patches.shape[1]}')
print(f'Elements per patch: {patches.shape[-1]}')

n = int(np.sqrt(patches.shape[1]))
plt.figure(figsize=(4, 4))

for i, patch in enumerate(patches[0]):
    ax = plt.subplot(n, n, i + 1)
    patch_img = tf.reshape(patch, (patch_size, patch_size, 3))
    plt.imshow(patch_img.numpy().astype('uint8'))
    plt.axis('off')
```

```
Image size: 224 X 224
Patch size: 7 X 7
Patches per image: 1024
Elements per patch: 147
```



ViT Model Architecture

```
model = tf.keras.Sequential([
    vit_model,
    tf.keras.layers.Flatten(),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Dense(128, activation = tf.nn.gelu),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Dense(64, activation = tf.nn.gelu),
    tf.keras.layers.Dense(32, activation = tf.nn.gelu),
    tf.keras.layers.Dense(4, 'softmax')
],
```



```
name = 'vision_transformer')
```

```
model.summary()
```

```
Model: "vision_transformer"
```

Layer (type)	Output Shape	Param #
vit-b16 (Functional)	(None, 768)	85798656
flatten_1 (Flatten)	(None, 768)	0
batch_normalization_2 (Batch Normalization)	(None, 768)	3072
dense_4 (Dense)	(None, 128)	98432
batch_normalization_3 (Batch Normalization)	(None, 128)	512
dense_5 (Dense)	(None, 64)	8256
dense_6 (Dense)	(None, 32)	2080
dense_7 (Dense)	(None, 4)	132
Total params: 85,911,140		
Trainable params: 85,909,348		
Non-trainable params: 1,792		

Training the Model

```
warnings.filterwarnings("ignore")
```

```
learning_rate = 1e-4
```

```
optimizer = tf.keras.optimizers.RectifiedAdam(learning_rate =  
learning_rate)
```

```
model.compile(optimizer = optimizer,  
              loss =  
tf.keras.losses.CategoricalCrossentropy(label_smoothing = 0.2),  
              metrics = ['accuracy'])
```

```
STEP_SIZE_TRAIN = train_gen.n // train_gen.batch_size  
STEP_SIZE_VALID = valid_gen.n // valid_gen.batch_size
```

```
early_stopping_callbacks = tf.keras.callbacks.EarlyStopping(patience =  
15, restore_best_weights = True, verbose = 1)
```



```
model.fit(x = train_gen,  
          steps_per_epoch = STEP_SIZE_TRAIN,  
          validation_data = valid_gen,  
          validation_steps = STEP_SIZE_VALID,  
          epochs = EPOCHS,  
          callbacks = early_stopping_callbacks)
```

Epoch 1/30

115/115 [=====] - 128s 756ms/step - loss:
1.2583 - accuracy: 0.5095 - val_loss: 1.1746 - val_accuracy: 0.5603

Epoch 2/30

115/115 [=====] - 70s 608ms/step - loss:
1.0260 - accuracy: 0.7009 - val_loss: 1.1565 - val_accuracy: 0.5603

Epoch 3/30

115/115 [=====] - 70s 606ms/step - loss:
0.9324 - accuracy: 0.7781 - val_loss: 1.1931 - val_accuracy: 0.5804

Epoch 4/30

115/115 [=====] - 70s 605ms/step - loss:
0.8840 - accuracy: 0.8140 - val_loss: 1.0552 - val_accuracy: 0.6763

Epoch 5/30

115/115 [=====] - 70s 612ms/step - loss:
0.8508 - accuracy: 0.8287 - val_loss: 1.0878 - val_accuracy: 0.6607

Epoch 6/30

115/115 [=====] - 70s 605ms/step - loss:
0.8329 - accuracy: 0.8407 - val_loss: 0.9854 - val_accuracy: 0.7455

Epoch 7/30

115/115 [=====] - 70s 608ms/step - loss:
0.8041 - accuracy: 0.8673 - val_loss: 0.9169 - val_accuracy: 0.7835

Epoch 8/30

115/115 [=====] - 70s 606ms/step - loss:
0.7928 - accuracy: 0.8635 - val_loss: 0.9503 - val_accuracy: 0.7545

Epoch 9/30

115/115 [=====] - 70s 608ms/step - loss:
0.7742 - accuracy: 0.8755 - val_loss: 0.9459 - val_accuracy: 0.7612

Epoch 10/30

115/115 [=====] - 69s 602ms/step - loss:
0.7581 - accuracy: 0.8880 - val_loss: 1.0175 - val_accuracy: 0.7299

Epoch 11/30

115/115 [=====] - 70s 604ms/step - loss:
0.7591 - accuracy: 0.8907 - val_loss: 0.9173 - val_accuracy: 0.7812

Epoch 12/30

115/115 [=====] - 70s 606ms/step - loss:
0.7466 - accuracy: 0.8978 - val_loss: 0.9606 - val_accuracy: 0.7746

Epoch 13/30

115/115 [=====] - 70s 611ms/step - loss:
0.7259 - accuracy: 0.9125 - val_loss: 0.9557 - val_accuracy: 0.7567

Epoch 14/30

115/115 [=====] - 70s 609ms/step - loss:
0.7141 - accuracy: 0.9233 - val_loss: 0.9600 - val_accuracy: 0.7589

Epoch 15/30
115/115 [=====] - 70s 607ms/step - loss:
0.7379 - accuracy: 0.8972 - val_loss: 0.9464 - val_accuracy: 0.7545
Epoch 16/30
115/115 [=====] - 70s 607ms/step - loss:
0.7098 - accuracy: 0.9233 - val_loss: 0.9352 - val_accuracy: 0.7790
Epoch 17/30
115/115 [=====] - 70s 606ms/step - loss:
0.7079 - accuracy: 0.9239 - val_loss: 0.8557 - val_accuracy: 0.8527
Epoch 18/30
115/115 [=====] - 70s 611ms/step - loss:
0.7025 - accuracy: 0.9282 - val_loss: 0.8681 - val_accuracy: 0.8103
Epoch 19/30
115/115 [=====] - 70s 608ms/step - loss:
0.6855 - accuracy: 0.9456 - val_loss: 0.8968 - val_accuracy: 0.8036
Epoch 20/30
115/115 [=====] - 70s 608ms/step - loss:
0.6887 - accuracy: 0.9380 - val_loss: 0.8516 - val_accuracy: 0.8237
Epoch 21/30
115/115 [=====] - 70s 606ms/step - loss:
0.6756 - accuracy: 0.9516 - val_loss: 0.9013 - val_accuracy: 0.8170
Epoch 22/30
115/115 [=====] - 70s 611ms/step - loss:
0.6692 - accuracy: 0.9565 - val_loss: 0.9370 - val_accuracy: 0.8036
Epoch 23/30
115/115 [=====] - 70s 606ms/step - loss:
0.6760 - accuracy: 0.9489 - val_loss: 0.8954 - val_accuracy: 0.8103
Epoch 24/30
115/115 [=====] - 70s 606ms/step - loss:
0.6772 - accuracy: 0.9511 - val_loss: 0.9165 - val_accuracy: 0.7991
Epoch 25/30
115/115 [=====] - 70s 605ms/step - loss:
0.6718 - accuracy: 0.9511 - val_loss: 0.9096 - val_accuracy: 0.8214
Epoch 26/30
115/115 [=====] - 70s 610ms/step - loss:
0.6802 - accuracy: 0.9424 - val_loss: 0.9591 - val_accuracy: 0.7679
Epoch 27/30
115/115 [=====] - 70s 605ms/step - loss:
0.6620 - accuracy: 0.9608 - val_loss: 0.9634 - val_accuracy: 0.7612
Epoch 28/30
115/115 [=====] - 70s 607ms/step - loss:
0.6813 - accuracy: 0.9451 - val_loss: 0.9122 - val_accuracy: 0.8013
Epoch 29/30
115/115 [=====] - 69s 602ms/step - loss:
0.6687 - accuracy: 0.9521 - val_loss: 1.0262 - val_accuracy: 0.7254
Epoch 30/30
115/115 [=====] - 70s 608ms/step - loss:
0.6601 - accuracy: 0.9581 - val_loss: 0.8990 - val_accuracy: 0.8103

<keras.callbacks.History at 0x7aa797b0af50>

```
# Save The Model
```

```
model.save('ViT_model.h5')
```


ViT Model Result

```
predicted_classes = np.argmax(model.predict(valid_gen, steps =
valid_gen.n // valid_gen.batch_size + 1), axis = 1)
true_classes = valid_gen.classes
class_labels = list(valid_gen.class_indices.keys())

confusionmatrix = confusion_matrix(true_classes, predicted_classes)
plt.figure(figsize = (16, 16))
sns.heatmap(confusionmatrix, cmap = 'Blues', annot = True, cbar =
True)

print(classification_report(true_classes, predicted_classes))
```

	precision	recall	f1-score	support
0	0.82	0.82	0.82	121
1	0.84	0.87	0.85	124
2	0.68	0.65	0.66	124
3	0.97	0.96	0.96	92
accuracy			0.82	461
macro avg	0.82	0.82	0.82	461
weighted avg	0.81	0.82	0.81	461



```
model.evaluate(valid_gen)
```

```
29/29 [=====] - 5s 156ms/step - loss: 0.8906  
- accuracy: 0.8156
```

```
[0.8906439542770386, 0.8156182169914246]
```

2 - ResNet50

```
model2 = tf.keras.Sequential()

base1= tf.keras.applications.ResNet50V2(include_top=False,
                                         input_shape=(224,224,3),
                                         pooling='avg',classes=4,
                                         weights='imagenet')

model2.add(base1)

model2.add(tf.keras.layers.Flatten())
model2.add(tf.keras.layers.BatchNormalization())
model2.add(tf.keras.layers.Dense(128, activation =
tfa.activations.gelu))
model2.add(tf.keras.layers.BatchNormalization())
model2.add(tf.keras.layers.Dense(64, activation =
tfa.activations.gelu))
model2.add(tf.keras.layers.Dropout(0.5))
model2.add(tf.keras.layers.Dense(32, activation =
tfa.activations.gelu))
model2.add(tf.keras.layers.Dense(4, 'softmax'))

for layer in base1.layers:
    layer.trainable = False

model2.summary()

Model: "sequential_4"
```

Layer (type)	Output Shape	Param #
resnet50v2 (Functional)	(None, 2048)	23564800
flatten_6 (Flatten)	(None, 2048)	0
batch_normalization_12 (Batch Normalization)	(None, 2048)	8192
dense_24 (Dense)	(None, 128)	262272
batch_normalization_13 (Batch Normalization)	(None, 128)	512
dense_25 (Dense)	(None, 64)	8256
dropout (Dropout)	(None, 64)	0
dense_26 (Dense)	(None, 32)	2080
dense_27 (Dense)	(None, 4)	132

Total params: 23,846,244
Trainable params: 277,092
Non-trainable params: 23,569,152

Training The Model

```
model2.compile(optimizer = optimizer,  
               loss =  
tf.keras.losses.CategoricalCrossentropy(label_smoothing = 0.2),  
               metrics = ['accuracy'])  
  
history = model2.fit(x = train_gen,  
                    steps_per_epoch = STEP_SIZE_TRAIN,  
                    validation_data = valid_gen,  
                    validation_steps = STEP_SIZE_VALID,  
                    epochs = EPOCHS,  
                    callbacks = early_stopping_callbacks)
```

Epoch 1/30
115/115 [=====] - 36s 281ms/step - loss:
1.2918 - accuracy: 0.5084 - val_loss: 1.1641 - val_accuracy: 0.6362
Epoch 2/30
115/115 [=====] - 31s 271ms/step - loss:
1.1255 - accuracy: 0.6161 - val_loss: 1.0899 - val_accuracy: 0.6942
Epoch 3/30
115/115 [=====] - 32s 275ms/step - loss:
1.0732 - accuracy: 0.6688 - val_loss: 1.0489 - val_accuracy: 0.7254
Epoch 4/30
115/115 [=====] - 32s 278ms/step - loss:
1.0472 - accuracy: 0.6922 - val_loss: 1.0374 - val_accuracy: 0.7254
Epoch 5/30
115/115 [=====] - 32s 276ms/step - loss:
1.0391 - accuracy: 0.6939 - val_loss: 1.0098 - val_accuracy: 0.7344
Epoch 6/30
115/115 [=====] - 32s 276ms/step - loss:
1.0299 - accuracy: 0.7031 - val_loss: 0.9939 - val_accuracy: 0.7478
Epoch 7/30
115/115 [=====] - 31s 271ms/step - loss:
1.0121 - accuracy: 0.7145 - val_loss: 0.9845 - val_accuracy: 0.7500
Epoch 8/30
115/115 [=====] - 31s 272ms/step - loss:
0.9870 - accuracy: 0.7401 - val_loss: 0.9793 - val_accuracy: 0.7522
Epoch 9/30
115/115 [=====] - 32s 274ms/step - loss:
0.9893 - accuracy: 0.7384 - val_loss: 0.9648 - val_accuracy: 0.7433
Epoch 10/30
115/115 [=====] - 32s 276ms/step - loss:
0.9786 - accuracy: 0.7395 - val_loss: 0.9772 - val_accuracy: 0.7455

Epoch 11/30
115/115 [=====] - 32s 275ms/step - loss:
0.9850 - accuracy: 0.7488 - val_loss: 0.9681 - val_accuracy: 0.7656
Epoch 12/30
115/115 [=====] - 31s 273ms/step - loss:
0.9552 - accuracy: 0.7597 - val_loss: 0.9575 - val_accuracy: 0.7746
Epoch 13/30
115/115 [=====] - 32s 275ms/step - loss:
0.9566 - accuracy: 0.7645 - val_loss: 0.9555 - val_accuracy: 0.7857
Epoch 14/30
115/115 [=====] - 31s 273ms/step - loss:
0.9411 - accuracy: 0.7711 - val_loss: 0.9462 - val_accuracy: 0.7768
Epoch 15/30
115/115 [=====] - 32s 274ms/step - loss:
0.9472 - accuracy: 0.7700 - val_loss: 0.9505 - val_accuracy: 0.7835
Epoch 16/30
115/115 [=====] - 31s 273ms/step - loss:
0.9435 - accuracy: 0.7678 - val_loss: 0.9492 - val_accuracy: 0.7768
Epoch 17/30
115/115 [=====] - 31s 272ms/step - loss:
0.9391 - accuracy: 0.7678 - val_loss: 0.9532 - val_accuracy: 0.7879
Epoch 18/30
115/115 [=====] - 32s 274ms/step - loss:
0.9380 - accuracy: 0.7760 - val_loss: 0.9452 - val_accuracy: 0.7857
Epoch 19/30
115/115 [=====] - 31s 273ms/step - loss:
0.9204 - accuracy: 0.7781 - val_loss: 0.9412 - val_accuracy: 0.7835
Epoch 20/30
115/115 [=====] - 31s 273ms/step - loss:
0.9211 - accuracy: 0.7847 - val_loss: 0.9382 - val_accuracy: 0.7857
Epoch 21/30
115/115 [=====] - 31s 273ms/step - loss:
0.9156 - accuracy: 0.7847 - val_loss: 0.9396 - val_accuracy: 0.7768
Epoch 22/30
115/115 [=====] - 31s 270ms/step - loss:
0.9006 - accuracy: 0.7983 - val_loss: 0.9377 - val_accuracy: 0.7924
Epoch 23/30
115/115 [=====] - 31s 273ms/step - loss:
0.9007 - accuracy: 0.7966 - val_loss: 0.9245 - val_accuracy: 0.7924
Epoch 24/30
115/115 [=====] - 31s 272ms/step - loss:
0.9115 - accuracy: 0.7852 - val_loss: 0.9207 - val_accuracy: 0.7879
Epoch 25/30
115/115 [=====] - 31s 270ms/step - loss:
0.9077 - accuracy: 0.7993 - val_loss: 0.9254 - val_accuracy: 0.7812
Epoch 26/30
115/115 [=====] - 31s 272ms/step - loss:
0.8963 - accuracy: 0.7950 - val_loss: 0.9267 - val_accuracy: 0.7969
Epoch 27/30




```
115/115 [=====] - 31s 270ms/step - loss:
0.8978 - accuracy: 0.8026 - val_loss: 0.9228 - val_accuracy: 0.7879
Epoch 28/30
115/115 [=====] - 31s 270ms/step - loss:
0.8913 - accuracy: 0.8227 - val_loss: 0.9170 - val_accuracy: 0.7991
Epoch 29/30
115/115 [=====] - 32s 275ms/step - loss:
0.8892 - accuracy: 0.8075 - val_loss: 0.9139 - val_accuracy: 0.8013
Epoch 30/30
115/115 [=====] - 31s 271ms/step - loss:
0.8953 - accuracy: 0.8032 - val_loss: 0.9062 - val_accuracy: 0.8058
```

Model Result

```
import matplotlib.pyplot as plt

# Plot training and validation accuracy
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

# Plot training and validation loss
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



```


predicted_classes = np.argmax(model2.predict(valid_gen, steps =
valid_gen.n // valid_gen.batch_size + 1), axis = 1)
true_classes = valid_gen.classes
class_labels = list(valid_gen.class_indices.keys())

```

```
confusionmatrix = confusion_matrix(true_classes, predicted_classes)
plt.figure(figsize = (16, 16))
sns.heatmap(confusionmatrix, cmap = 'Blues', annot = True, cbar =
True)
```

```
print(classification_report(true_classes, predicted_classes))
```

	precision	recall	f1-score	support
0	0.77	0.79	0.78	121
1	0.87	0.94	0.90	124
2	0.70	0.59	0.64	124
3	0.89	0.93	0.91	92
accuracy			0.81	461
macro avg	0.81	0.82	0.81	461
weighted avg	0.80	0.81	0.80	461



3 - VGG19

```
model3 = tf.keras.Sequential()  
base2= tf.keras.applications.VGG19(include_top=False,  
    input_shape=(224,224,3),  
    pooling='avg',classes=4,  
    weights='imagenet')
```

```

model3.add(base2)

model3.add(tf.keras.layers.Flatten())
model3.add(tf.keras.layers.BatchNormalization())
model3.add(tf.keras.layers.Dense(128, activation =
tfa.activations.gelu))
model3.add(tf.keras.layers.BatchNormalization())
model3.add(tf.keras.layers.Dense(64, activation =
tfa.activations.gelu))
model3.add(tf.keras.layers.Dropout(0.5))
model3.add(tf.keras.layers.Dense(32, activation =
tfa.activations.gelu))
model3.add(tf.keras.layers.Dense(4, 'softmax'))

for layer in base2.layers:
    layer.trainable = False

```

```
model3.summary()
```

```
Model: "sequential_5"
```

Layer (type)	Output Shape	Param #
vgg19 (Functional)	(None, 512)	20024384
flatten_7 (Flatten)	(None, 512)	0
batch_normalization_14 (Batch Normalization)	(None, 512)	2048
dense_28 (Dense)	(None, 128)	65664
batch_normalization_15 (Batch Normalization)	(None, 128)	512
dense_29 (Dense)	(None, 64)	8256
dropout_1 (Dropout)	(None, 64)	0
dense_30 (Dense)	(None, 32)	2080
dense_31 (Dense)	(None, 4)	132
Total params: 20,103,076		
Trainable params: 77,412		
Non-trainable params: 20,025,664		

Training The Model

```

model3.compile(optimizer = optimizer,
               loss =

```

```

tf.keras.losses.CategoricalCrossentropy(label_smoothing = 0.2),
    metrics = ['accuracy'])

history = model3.fit(x = train_gen,
    steps_per_epoch = STEP_SIZE_TRAIN,
    validation_data = valid_gen,
    validation_steps = STEP_SIZE_VALID,
    epochs = EPOCHS,
    callbacks = early_stopping_callbacks)

Epoch 1/30
115/115 [=====] - 37s 298ms/step - loss:
1.2975 - accuracy: 0.4753 - val_loss: 1.3571 - val_accuracy: 0.3214
Epoch 2/30
115/115 [=====] - 34s 294ms/step - loss:
1.1385 - accuracy: 0.6063 - val_loss: 1.3191 - val_accuracy: 0.3951
Epoch 3/30
115/115 [=====] - 34s 292ms/step - loss:
1.0992 - accuracy: 0.6580 - val_loss: 1.2927 - val_accuracy: 0.4509
Epoch 4/30
115/115 [=====] - 33s 289ms/step - loss:
1.0624 - accuracy: 0.6824 - val_loss: 1.3154 - val_accuracy: 0.3929
Epoch 5/30
115/115 [=====] - 33s 288ms/step - loss:
1.0470 - accuracy: 0.6977 - val_loss: 1.3571 - val_accuracy: 0.3750
Epoch 6/30
115/115 [=====] - 33s 291ms/step - loss:
1.0485 - accuracy: 0.6813 - val_loss: 1.3676 - val_accuracy: 0.4129
Epoch 7/30
115/115 [=====] - 34s 292ms/step - loss:
1.0282 - accuracy: 0.7085 - val_loss: 1.3196 - val_accuracy: 0.4531
Epoch 8/30
115/115 [=====] - 33s 287ms/step - loss:
1.0243 - accuracy: 0.7020 - val_loss: 1.3624 - val_accuracy: 0.4330
Epoch 9/30
115/115 [=====] - 33s 291ms/step - loss:
1.0101 - accuracy: 0.7129 - val_loss: 1.3739 - val_accuracy: 0.4308
Epoch 10/30
115/115 [=====] - 33s 289ms/step - loss:
1.0001 - accuracy: 0.7259 - val_loss: 1.3860 - val_accuracy: 0.4196
Epoch 11/30
115/115 [=====] - 34s 292ms/step - loss:
0.9917 - accuracy: 0.7346 - val_loss: 1.3685 - val_accuracy: 0.4308
Epoch 12/30
115/115 [=====] - 33s 287ms/step - loss:
1.0046 - accuracy: 0.7134 - val_loss: 1.3531 - val_accuracy: 0.4397
Epoch 13/30
115/115 [=====] - 34s 291ms/step - loss:
0.9943 - accuracy: 0.7319 - val_loss: 1.3407 - val_accuracy: 0.4554
Epoch 14/30

```



```
115/115 [=====] - 33s 290ms/step - loss:
0.9756 - accuracy: 0.7395 - val_loss: 1.3269 - val_accuracy: 0.4643
Epoch 15/30
115/115 [=====] - 34s 293ms/step - loss:
0.9797 - accuracy: 0.7412 - val_loss: 1.3535 - val_accuracy: 0.4576
Epoch 16/30
115/115 [=====] - 33s 287ms/step - loss:
0.9696 - accuracy: 0.7482 - val_loss: 1.3601 - val_accuracy: 0.4375
Epoch 17/30
115/115 [=====] - 34s 292ms/step - loss:
0.9656 - accuracy: 0.7520 - val_loss: 1.3644 - val_accuracy: 0.4442
Epoch 18/30
115/115 [=====] - 33s 289ms/step - loss:
0.9634 - accuracy: 0.7482 - val_loss: 1.3458 - val_accuracy: 0.4598
Restoring model weights from the end of the best epoch.
Epoch 00018: early stopping
```

Model Result

```
import matplotlib.pyplot as plt

# Plot training and validation accuracy
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

# Plot training and validation loss
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

4 - InceptionV3

```
model4 = tf.keras.Sequential()

base3=
tf.keras.applications.inception_v3.InceptionV3(include_top=False,
        input_shape=(224,224,3),
        pooling='avg',
        classes=4,
        weights='imagenet')

model4.add(base3)

model4.add(tf.keras.layers.Flatten())
model4.add(tf.keras.layers.BatchNormalization())
model4.add(tf.keras.layers.Dense(128, activation =
tfa.activations.gelu))
model4.add(tf.keras.layers.BatchNormalization())
model4.add(tf.keras.layers.Dense(64, activation =
tfa.activations.gelu))
model4.add(tf.keras.layers.Dropout(0.5))
model4.add(tf.keras.layers.Dense(32, activation =
tfa.activations.gelu))
model4.add(tf.keras.layers.Dense(4, 'softmax'))

for layer in base3.layers:
    layer.trainable = False

model4.summary()

Model: "sequential_7"
```

Layer (type)	Output Shape	Param #
inception_v3 (Functional)	(None, 2048)	21802784
flatten_9 (Flatten)	(None, 2048)	0
batch_normalization_206 (Bat	(None, 2048)	8192
dense_36 (Dense)	(None, 128)	262272
batch_normalization_207 (Bat	(None, 128)	512
dense_37 (Dense)	(None, 64)	8256
dropout_3 (Dropout)	(None, 64)	0
dense_38 (Dense)	(None, 32)	2080

dense_39 (Dense)	(None, 4)	132
------------------	-----------	-----

Total params: 22,084,228
Trainable params: 277,092
Non-trainable params: 21,807,136

Training The Model

```
model4.compile(optimizer = optimizer,  
               loss =  
tf.keras.losses.CategoricalCrossentropy(label_smoothing = 0.2),  
               metrics = ['accuracy'])  
  
history = model4.fit(x = train_gen,  
                    steps_per_epoch = STEP_SIZE_TRAIN,  
                    validation_data = valid_gen,  
                    validation_steps = STEP_SIZE_VALID,  
                    epochs = EPOCHS,  
                    callbacks = early_stopping_callbacks)
```

Epoch 1/30
115/115 [=====] - 37s 276ms/step - loss:
1.2980 - accuracy: 0.5128 - val_loss: 1.1558 - val_accuracy: 0.5893
Epoch 2/30
115/115 [=====] - 31s 268ms/step - loss:
1.1102 - accuracy: 0.6291 - val_loss: 1.0686 - val_accuracy: 0.6830
Epoch 3/30
115/115 [=====] - 31s 273ms/step - loss:
1.0781 - accuracy: 0.6591 - val_loss: 1.0270 - val_accuracy: 0.6964
Epoch 4/30
115/115 [=====] - 31s 267ms/step - loss:
1.0635 - accuracy: 0.6688 - val_loss: 1.0031 - val_accuracy: 0.7076
Epoch 5/30
115/115 [=====] - 31s 269ms/step - loss:
1.0283 - accuracy: 0.7031 - val_loss: 0.9739 - val_accuracy: 0.7321
Epoch 6/30
115/115 [=====] - 31s 266ms/step - loss:
1.0143 - accuracy: 0.7151 - val_loss: 0.9812 - val_accuracy: 0.7321
Epoch 7/30
115/115 [=====] - 31s 268ms/step - loss:
1.0105 - accuracy: 0.7069 - val_loss: 0.9713 - val_accuracy: 0.7411
Epoch 8/30
115/115 [=====] - 30s 265ms/step - loss:
0.9875 - accuracy: 0.7319 - val_loss: 0.9604 - val_accuracy: 0.7277
Epoch 9/30
115/115 [=====] - 31s 268ms/step - loss:
0.9854 - accuracy: 0.7178 - val_loss: 0.9477 - val_accuracy: 0.7589
Epoch 10/30

```
115/115 [=====] - 30s 264ms/step - loss:
0.9728 - accuracy: 0.7428 - val_loss: 0.9441 - val_accuracy: 0.7701
Epoch 11/30
115/115 [=====] - 31s 267ms/step - loss:
0.9871 - accuracy: 0.7299 - val_loss: 0.9465 - val_accuracy: 0.7656
Epoch 12/30
115/115 [=====] - 31s 267ms/step - loss:
0.9553 - accuracy: 0.7607 - val_loss: 0.9393 - val_accuracy: 0.7701
Epoch 13/30
115/115 [=====] - 30s 263ms/step - loss:
0.9599 - accuracy: 0.7531 - val_loss: 0.9306 - val_accuracy: 0.7701
Epoch 14/30
115/115 [=====] - 30s 263ms/step - loss:
0.9593 - accuracy: 0.7402 - val_loss: 0.9406 - val_accuracy: 0.7723
Epoch 15/30
115/115 [=====] - 30s 264ms/step - loss:
0.9477 - accuracy: 0.7597 - val_loss: 0.9360 - val_accuracy: 0.7656
Epoch 16/30
115/115 [=====] - 31s 266ms/step - loss:
0.9381 - accuracy: 0.7548 - val_loss: 0.9161 - val_accuracy: 0.7857
Epoch 17/30
115/115 [=====] - 30s 262ms/step - loss:
0.9431 - accuracy: 0.7673 - val_loss: 0.9177 - val_accuracy: 0.7902
Epoch 18/30
115/115 [=====] - 31s 266ms/step - loss:
0.9262 - accuracy: 0.7743 - val_loss: 0.9170 - val_accuracy: 0.7902
Epoch 19/30
115/115 [=====] - 30s 265ms/step - loss:
0.9427 - accuracy: 0.7694 - val_loss: 0.9112 - val_accuracy: 0.7924
Epoch 20/30
115/115 [=====] - 31s 265ms/step - loss:
0.9366 - accuracy: 0.7760 - val_loss: 0.9107 - val_accuracy: 0.7924
Epoch 21/30
115/115 [=====] - 31s 266ms/step - loss:
0.9209 - accuracy: 0.7879 - val_loss: 0.9193 - val_accuracy: 0.7946
Epoch 22/30
115/115 [=====] - 31s 267ms/step - loss:
0.9274 - accuracy: 0.7771 - val_loss: 0.9146 - val_accuracy: 0.7902
Epoch 23/30
115/115 [=====] - 30s 264ms/step - loss:
0.9111 - accuracy: 0.7923 - val_loss: 0.9135 - val_accuracy: 0.7857
Epoch 24/30
115/115 [=====] - 31s 266ms/step - loss:
0.9050 - accuracy: 0.7901 - val_loss: 0.9098 - val_accuracy: 0.7790
Epoch 25/30
115/115 [=====] - 30s 262ms/step - loss:
0.9179 - accuracy: 0.7852 - val_loss: 0.9059 - val_accuracy: 0.7946
Epoch 26/30
115/115 [=====] - 30s 263ms/step - loss:
```


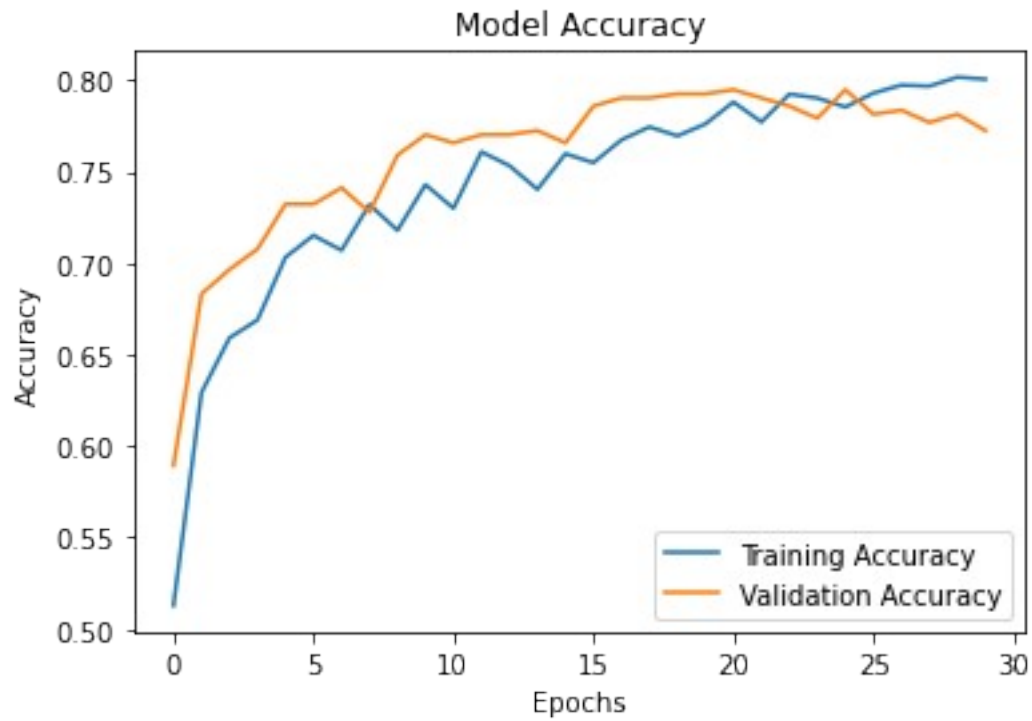
```
0.9064 - accuracy: 0.7928 - val_loss: 0.9140 - val_accuracy: 0.7812
Epoch 27/30
115/115 [=====] - 30s 262ms/step - loss:
0.9064 - accuracy: 0.7972 - val_loss: 0.9173 - val_accuracy: 0.7835
Epoch 28/30
115/115 [=====] - 30s 263ms/step - loss:
0.9056 - accuracy: 0.7966 - val_loss: 0.9181 - val_accuracy: 0.7768
Epoch 29/30
115/115 [=====] - 30s 265ms/step - loss:
0.8980 - accuracy: 0.8015 - val_loss: 0.9099 - val_accuracy: 0.7812
Epoch 30/30
115/115 [=====] - 30s 265ms/step - loss:
0.8901 - accuracy: 0.8004 - val_loss: 0.9140 - val_accuracy: 0.7723
```

Model Result

```
import matplotlib.pyplot as plt

# Plot training and validation accuracy
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

# Plot training and validation loss
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```




```
predicted_classes = np.argmax(model4.predict(valid_gen, steps =  
valid_gen.n // valid_gen.batch_size + 1), axis = 1)  
true_classes = valid_gen.classes  
class_labels = list(valid_gen.class_indices.keys())
```

```
confusionmatrix = confusion_matrix(true_classes, predicted_classes)
plt.figure(figsize = (16, 16))
sns.heatmap(confusionmatrix, cmap = 'Blues', annot = True, cbar =
True)
```

```
print(classification_report(true_classes, predicted_classes))
```

	precision	recall	f1-score	support
0	0.70	0.79	0.74	121
1	0.96	0.77	0.85	124
2	0.62	0.63	0.63	124
3	0.90	0.99	0.94	92
accuracy			0.78	461
macro avg	0.80	0.79	0.79	461
weighted avg	0.79	0.78	0.78	461



```
model2.save("resnet_model.h5")  
model3.save("vgg19.h5")  
model4.save("inception_model.h5")
```

5 - EfficientNet-B0

```
model5 = tf.keras.Sequential()

base5= tf.keras.applications.EfficientNetB0(include_top=False,
      input_shape=(224,224,3),
      pooling='avg',
      classes=4,
      weights='imagenet')

model5.add(base5)

model5.add(tf.keras.layers.Flatten())
model5.add(tf.keras.layers.BatchNormalization())
model5.add(tf.keras.layers.Dense(128, activation = "relu"))
model5.add(tf.keras.layers.BatchNormalization())
model5.add(tf.keras.layers.Dense(64, activation = "relu"))
model5.add(tf.keras.layers.Dropout(0.5))
model5.add(tf.keras.layers.Dense(32, activation = "relu"))
model5.add(tf.keras.layers.Dense(4, 'softmax'))

for layer in base5.layers:
    layer.trainable = False
```

```
model4.summary()
```

```
Model: "sequential_7"
```

Layer (type)	Output Shape	Param #
inception_v3 (Functional)	(None, 2048)	21802784
flatten_9 (Flatten)	(None, 2048)	0
batch_normalization_206 (Batch Normalization)	(None, 2048)	8192
dense_36 (Dense)	(None, 128)	262272
batch_normalization_207 (Batch Normalization)	(None, 128)	512
dense_37 (Dense)	(None, 64)	8256
dropout_3 (Dropout)	(None, 64)	0
dense_38 (Dense)	(None, 32)	2080
dense_39 (Dense)	(None, 4)	132
Total params: 22,084,228		
Trainable params: 277,092		

Non-trainable params: 21,807,136

Training the model

```
model5.compile(optimizer = optimizer,
               loss =
tf.keras.losses.CategoricalCrossentropy(label_smoothing = 0.2),
               metrics = ['accuracy'])

history = model5.fit(x = train_gen,
                    steps_per_epoch = STEP_SIZE_TRAIN,
                    validation_data = valid_gen,
                    validation_steps = STEP_SIZE_VALID,
                    epochs = EPOCHS,
                    callbacks = early_stopping_callbacks)
```

Epoch 1/30
115/115 [=====] - 40s 296ms/step - loss:
1.4786 - accuracy: 0.2926 - val_loss: 1.3873 - val_accuracy: 0.2768

Epoch 2/30
115/115 [=====] - 30s 261ms/step - loss:
1.4124 - accuracy: 0.2964 - val_loss: 1.3918 - val_accuracy: 0.2768

Epoch 3/30
115/115 [=====] - 30s 262ms/step - loss:
1.4032 - accuracy: 0.3051 - val_loss: 1.3879 - val_accuracy: 0.2768

Epoch 4/30
115/115 [=====] - 30s 264ms/step - loss:
1.3893 - accuracy: 0.3002 - val_loss: 1.3778 - val_accuracy: 0.2768

Epoch 5/30
115/115 [=====] - 31s 268ms/step - loss:
1.3664 - accuracy: 0.3241 - val_loss: 1.3698 - val_accuracy: 0.3929

Epoch 6/30
115/115 [=====] - 31s 268ms/step - loss:
1.3646 - accuracy: 0.3235 - val_loss: 1.3742 - val_accuracy: 0.2701

Epoch 7/30
115/115 [=====] - 30s 262ms/step - loss:
1.3615 - accuracy: 0.3350 - val_loss: 1.3932 - val_accuracy: 0.2701

Epoch 8/30
115/115 [=====] - 30s 265ms/step - loss:
1.3557 - accuracy: 0.3361 - val_loss: 1.3811 - val_accuracy: 0.2701

Epoch 9/30
115/115 [=====] - 30s 264ms/step - loss:
1.3616 - accuracy: 0.3355 - val_loss: 1.4124 - val_accuracy: 0.2522

Epoch 10/30
115/115 [=====] - 30s 261ms/step - loss:
1.3611 - accuracy: 0.3143 - val_loss: 1.3773 - val_accuracy: 0.2701

Epoch 11/30
115/115 [=====] - 30s 260ms/step - loss:

1.3551 - accuracy: 0.3284 - val_loss: 1.3875 - val_accuracy: 0.3504
Epoch 12/30
115/115 [=====] - 30s 264ms/step - loss:
1.3642 - accuracy: 0.3056 - val_loss: 1.3984 - val_accuracy: 0.1942
Epoch 13/30
115/115 [=====] - 30s 260ms/step - loss:
1.3434 - accuracy: 0.3442 - val_loss: 1.4819 - val_accuracy: 0.1607
Epoch 14/30
115/115 [=====] - 30s 260ms/step - loss:
1.3450 - accuracy: 0.3469 - val_loss: 1.4188 - val_accuracy: 0.1763
Epoch 15/30
115/115 [=====] - 30s 260ms/step - loss:
1.3522 - accuracy: 0.3409 - val_loss: 1.5284 - val_accuracy: 0.1741
Epoch 16/30
115/115 [=====] - 30s 263ms/step - loss:
1.3430 - accuracy: 0.3709 - val_loss: 1.4088 - val_accuracy: 0.3438
Epoch 17/30
115/115 [=====] - 30s 261ms/step - loss:
1.3486 - accuracy: 0.3409 - val_loss: 1.3867 - val_accuracy: 0.1473
Epoch 18/30
115/115 [=====] - 30s 261ms/step - loss:
1.3530 - accuracy: 0.3301 - val_loss: 1.3672 - val_accuracy: 0.3080
Epoch 19/30
115/115 [=====] - 30s 263ms/step - loss:
1.3407 - accuracy: 0.3339 - val_loss: 1.3751 - val_accuracy: 0.2701
Epoch 20/30
115/115 [=====] - 31s 265ms/step - loss:
1.3446 - accuracy: 0.3350 - val_loss: 1.3740 - val_accuracy: 0.2969
Epoch 21/30
115/115 [=====] - 30s 262ms/step - loss:
1.3367 - accuracy: 0.3415 - val_loss: 1.4608 - val_accuracy: 0.1317
Epoch 22/30
115/115 [=====] - 31s 266ms/step - loss:
1.3382 - accuracy: 0.3535 - val_loss: 1.3900 - val_accuracy: 0.3281
Epoch 23/30
115/115 [=====] - 31s 268ms/step - loss:
1.3348 - accuracy: 0.3622 - val_loss: 1.4101 - val_accuracy: 0.2210
Epoch 24/30
115/115 [=====] - 31s 268ms/step - loss:
1.3362 - accuracy: 0.3448 - val_loss: 1.4264 - val_accuracy: 0.3460
Epoch 25/30
115/115 [=====] - 30s 264ms/step - loss:
1.3349 - accuracy: 0.3540 - val_loss: 1.4486 - val_accuracy: 0.3214
Epoch 26/30
115/115 [=====] - 31s 267ms/step - loss:
1.3290 - accuracy: 0.3768 - val_loss: 1.4613 - val_accuracy: 0.1808
Epoch 27/30
115/115 [=====] - 30s 265ms/step - loss:
1.3229 - accuracy: 0.3654 - val_loss: 1.4471 - val_accuracy: 0.1362

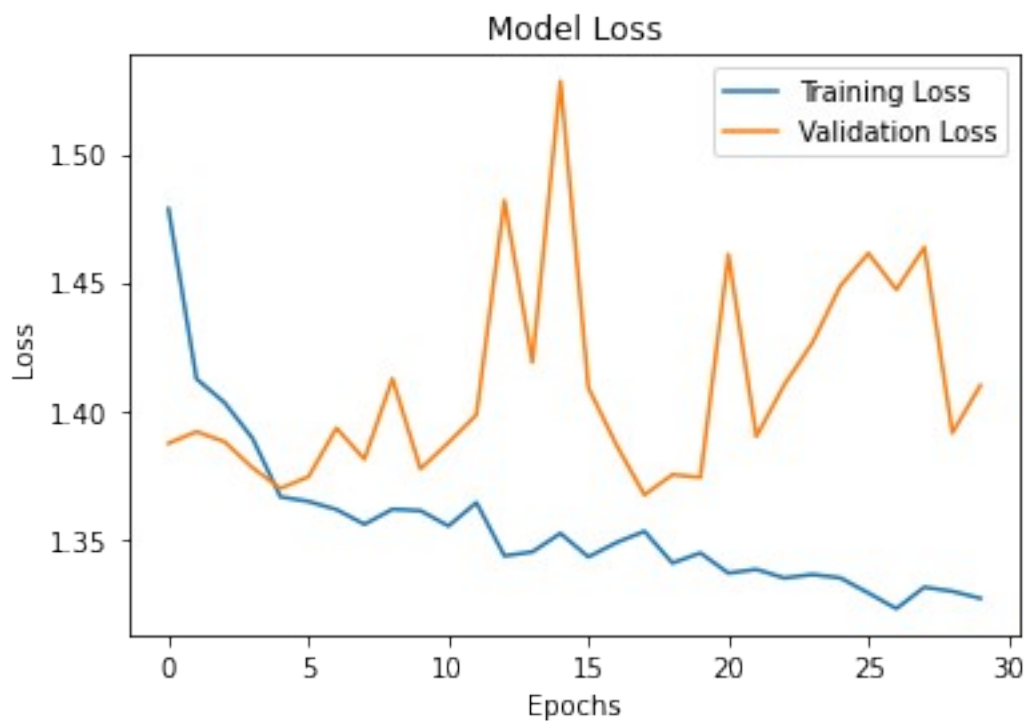

```
Epoch 28/30
115/115 [=====] - 31s 268ms/step - loss:
1.3313 - accuracy: 0.3589 - val_loss: 1.4636 - val_accuracy: 0.1339
Epoch 29/30
115/115 [=====] - 31s 272ms/step - loss:
1.3296 - accuracy: 0.3562 - val_loss: 1.3914 - val_accuracy: 0.3281
Epoch 30/30
115/115 [=====] - 31s 267ms/step - loss:
1.3269 - accuracy: 0.3741 - val_loss: 1.4098 - val_accuracy: 0.1674
```

Model Result

```
import matplotlib.pyplot as plt

# Plot training and validation accuracy
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

# Plot training and validation loss
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



```

predicted_classes = np.argmax(model5.predict(valid_gen, steps =
valid_gen.n // valid_gen.batch_size + 1), axis = 1)
true_classes = valid_gen.classes
class_labels = list(valid_gen.class_indices.keys())

```

```
confusionmatrix = confusion_matrix(true_classes, predicted_classes)
plt.figure(figsize = (16, 16))
sns.heatmap(confusionmatrix, cmap = 'Blues', annot = True, cbar =
True)
```

```
print(classification_report(true_classes, predicted_classes))
```

	precision	recall	f1-score	support
0	0.14	0.16	0.15	121
1	0.00	0.00	0.00	124
2	0.00	0.00	0.00	124
3	0.21	0.74	0.33	92
accuracy			0.19	461
macro avg	0.09	0.22	0.12	461
weighted avg	0.08	0.19	0.10	461

